

HOLLAND

ARCHITECTURE

- Broken into two logical components
 - holland-core (single package)
 - holland-plugins (multiple packages)

HOLLAND.CORE

- Scaffold/Framework for building backup interfaces
- No special knowledge of MySQL, LVM, Postgres, etc.
 - (These are all encapsulated in plugins)
- Currently has two frontends: “holland” and “holland_commvault”
 - (Hopefully more interesting interfaces in the future)

HOLLAND.CORE (CONTINUED)

- holland.core consists of 5 main components:
 - config, plugin, backup, command and backports
 - Also includes the “holland” script as an interface to various builtin-commands
 - (This will probably be refactored at some point)
 - Also includes some utility functions used in plugins and core

HOLLAND.CORE.CONFIG

- Main interface to `holland.conf` and `global-plugin/backupset` configurations
- Built on top of Michael Foord's `configobj`
- Handles merging global and per-backupset configs

HOLLAND.CORE.PLUGIN

- Loads plugins :)
 - (really, loads setuptools “entrypoints”)
- Uses (and currently married-to) `pkg_resources`
 - as is trac, python-nose, paster, etc.
- <http://ianbicking.org/docs/pycon2006/plugins.html>

HOLLAND.CORE.PLUGIN (CONTINUED)

- `add_plugin_dir(plugin_directory)`
- `load_first_entrypoint(group, pluginname)`
- `load_backup_plugin(name)`
 - = `load_first_entrypoint('holland.backup'.name)`
- `get_commands(), iter_plugininfo(), iter_dists()`

HOLLAND.CORE.BACKUP

- Sets up backup “spool” where backup data is stored
- Loads backupset configs
- Loads plugins
- Dispatches to backup plugins
- Understands “dry-run”
 - Plugins must implement dry-run separately :(

HOLLAND.CORE.BACKUP (CONTINUED)

- `backup(backupset_name, dry_run=False)`
- Saves a copy of backupset config as \${dst}/`backup.conf`
- Relies on `holland.core.spool.py` to manage the backup destination
- raises `PluginLoadError` or `BackupError` exceptions

HOLLAND.CORE.COMMAND

- Pluggable command support
- Mostly intended to support subcommands for “holland”
- Also exposes a general run(*args) API usable by external code
- Uses optparse

HOLLAND.CORE.COMMAND (CONTINUED)

```
class MyCommand(Command):

    """${cmd_usage}

    ${cmd_option_list}

    """

    name = 'foo'

    aliases = ['f']

    options = [ option('--foo-the-bar', action=store_true) ]

    description = 'foo a bar'

    def run(self, cmd, options, bar):

        ...
```

HOLLAND.CORE.COMMAND (CONTINUED)

- Command superclass actually calls run via dispatch (*args)
 - introspects arg names from run() signature via inspect
 - (recent argparse patches do this too)
- options[] just set of optparse options
 - (should probably be converted to argparse :)

HOLLAND.CORE.BACKPORTS

- Backports of standard python libraries
- mostly for python-2.3 support
- useful for making module behavior more consistent
(e.g. subprocess EINTR enhancements from debian)
- holland is currently naughty and loads these in
`holland.core.__init__` unconditionally

HOLLAND PLUGINS

- Plugins are packaged as python eggs
 - Holland refers to entrypoints as plugins
- Three types of plugins: backup, command and lib
 - (lib isn't really a “plugin”)
- Standard location /usr/share/holland/plugins
- Usually part of a holland namespace package

HOLLAND BACKUP PLUGINS

- Three basic methods:
 - `__init__(backupset_name, config, directory, dry_run)`
 - `estimate_backup_size()`
 - `backup()`
- `__init__` can be any callable that returns an object

HOLLAND.BACKUP.MYSQLDUMP

- Wraps ‘mysqldump’
- Does size estimations based on `information_schema`
- supports mysqldump lock modes (none, per-database, full-server, single-transaction)
 - also supports auto-detect for single-transaction
- Database and Table filtering via glob patterns

HOLLAND.COMMANDS

- 90% of commands are in `holland-core`
- Some proof-of-concept mysql tools in `plugins/holland.commands.mysql`
- Will probably extend this framework to provide export functionality via sub-sub-commands
 - `holland export backupset [backupset-exporter options]`

HOLLAND.LIB

- Two main packages here: `holland.lib.common`, `holland.lib.mysql`
- Just normal packages that provide functionality for plugins (never used by `holland.core` directly)
- Some of these could probably be useful outside of `holland`

HOLLAND.LIB.COMMON

- “Common” functionality
- currently compression and archiving
 - Also a multidict implementation from Ian Bicking
 - Also a portable ‘which’ implementation to find commands
 - Fairly simple

HOLLAND.LIB.COMMON - COMPRESSION

- Provides python file-like objects to wrap compression/decompression
- Always expose a usable `fileno()` attribute
 - Therefore, usable directly by `subprocess/fork`
- Methods (`gzip,bzip2,lzop`) hardcoded in dictionary :(
 - (Should be pluggable?)

HOLLAND.LIB.COMMON (MORE) COMPRESSION

- `CompressionOutput` = write out compressed files
 - supports inline flag, where it will compress on close()
- `CompressionInput` = read and uncompress files
- Based on subprocess, dispatching to a command
- Low python overhead, but portability issues (?)

HOLLAND.LIB.COMMON ARCHIVING

- Only used by `holland.backup.mysqlhotcopy`
 - (no one uses `holland.backup.mysqlhotcopy`)
- Intended to abstract zip, tar, or regular directory access
- zip, tar pure python, but perform well-enough
- probably need some cleanup

HOLLAND.LIB.MYSQL

- Used by MySQL plugins
 - Largely `holland.backup.mysqldump`
- Provides 4 main bits of functionality: client connections, `my.cnf` parsing, schema filtering and `mysql` command introspection
- Hard to test independently currently

HOLLAND.LIB.MYSQL.CLIENT

THE CLIENT OBJECT

- Wraps MySQLdb.Connection
 - Can be used as drop-in replacement for MySQLdb.Connection
- Provides useful common functionality used by DBAs on the ‘mysql’ cli tool
 - show_databases, show_tables, flush_*, show_slave_status, show_master_status, etc.

HOLLAND.LIB.MYSQL.FIND SCHEMA FILTERING

- Tightly bound to dbapi :(
- Inspired heavily by Maatkit's MySQLFind class
- Uses inclusion,exclusion lists for databases and tables
- Allows negative and positive searching
 - mysqldump only supports --ignore-table=db.tbl :(

HOLLAND.LIB.MYSQL.OPTION OPTION FILE PARSING

- <http://dev.mysql.com/doc/refman/5.1/en/option-files.html>
- Used to merge multiple .my.cnf files
 - Also to write out large mysqldump ignore-table lists for filtering
 - my.cnf may use shortest-unique-prefix matching :(
 - Other weird behavior of MySQL

HOLLAND.LIB.MYSQL.CLI COMMAND INTROSPECTION

- Derived from DBA scripts
- Mostly used to parse current version of utilities
 - `mysqldump.version < mysqlserver.version?`
 - Also makes it easy to check if a utility (or mysqld) supports some option
 - Although version checks are usually quite sufficient

THINGS TO DO

- Test and document
- Code cleanup
- Fix bugs
- LVM support in high demand (close to being ready)
- Export functionality requested (some good ideas here)

THINGS TO DO TESTING

- `holland.lib.common` is short and easy to test/document
 - Code is simple and should be easy to refactor
- `holland-core`
 - backup, spool needs cleanup and better defined behavior
 - `holland.commands` could be more robust

THINGS TO DO END-USER TESTING

- Try crazy stuff with holland backup
 - Stress the table/database filtering for mysqldump
 - Check backup and restore integrity
 - Performance testing - large schemas, compression
 - Find broken subcommands or non-intuitive behavior

THINGS TO DO DOCUMENTATION

- User documentation is pretty good from hollandaise 2.0 (Thanks Sweetums :)
- Developer documentation is non-existent (blame Andy)
- How should we document everything?
 - monolithic or manual-per-plugin?

THINGS TO DO NICE TO HAVES

- Signal support (clean behavior on SIGTERM,SIGHUP,etc.)
- Export functionality - some high level design on gforge
- Refactoring holland.lib.mysql.find to separate filtering from dbapi
- ???